

dplyr Fundamentals

Tidying and manipulating data

Download the section 5 .Rmd handout to
STAT240/lecture/sect05-dplyr.

Download two datasets to STAT240/data:

- penguins.csv
- madison-weather-1869-2023.csv

Material in this section is covered by Chapters 7-8
on the notes website.

dplyr is a collection of command for tidying data.

- Specific “grammar” just like ggplot.

Here is a cheatsheet of dplyr commands.

The **pipe** operator `%>%` takes an object and passes it into a command.

- Works like the word “then”
- Useful for a sequence of functions

 $\%>\%$ `mix()` $\%>\%$ `bake()``frost()`

Basic dplyr command:

- Input a df
- Alter the df
- Output the new df

Chain commands together with the pipe.

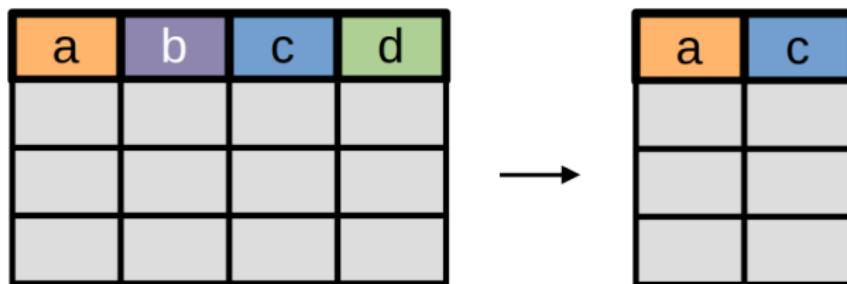
Column commands act on the columns and leave rows unchanged.

- `select()`, `relocate()`, `rename()`,
`mutate()`

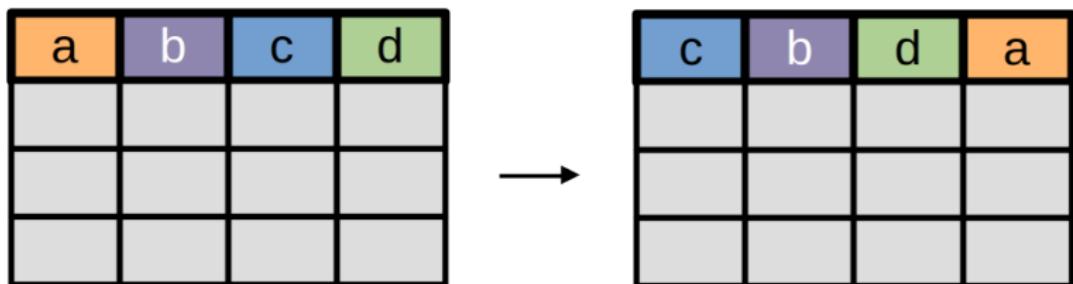
Rows commands act on the rows and leave columns unchanged.

- `arrange()`, `filter()`, `drop_na()`

`select()` reduces to specific columns.



`relocate()` moves columns.



`rename()` renames columns.

a	b	c	d



i	ii	iii	iv

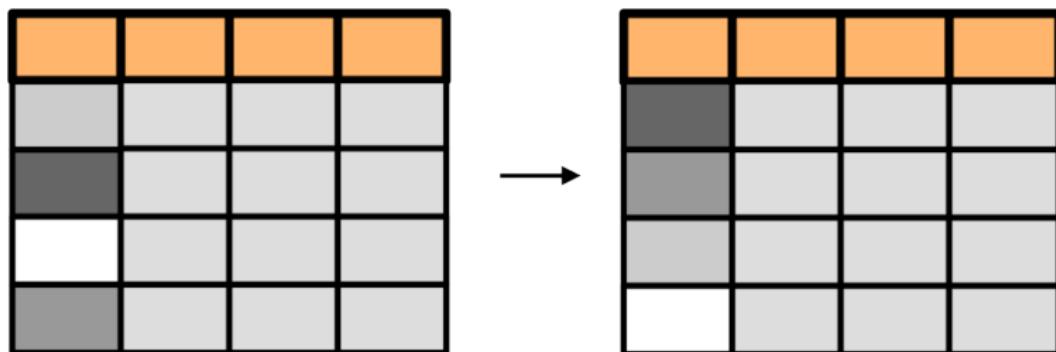
`mutate()` creates a new column.

a	b	c	d

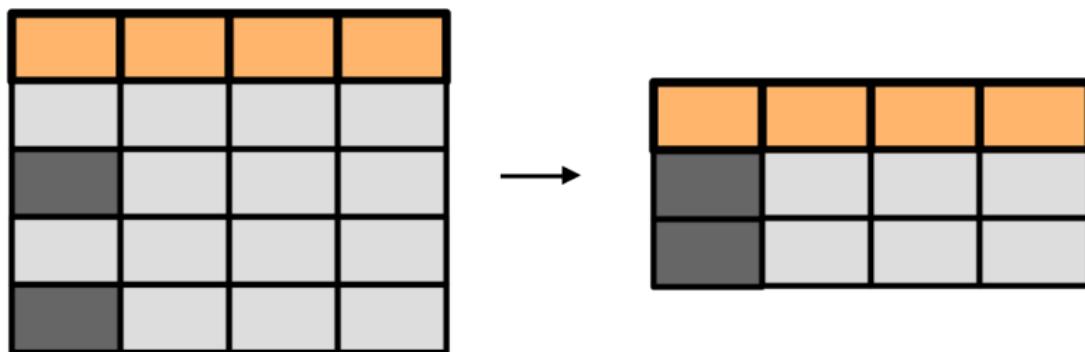


a	b	c	d	e

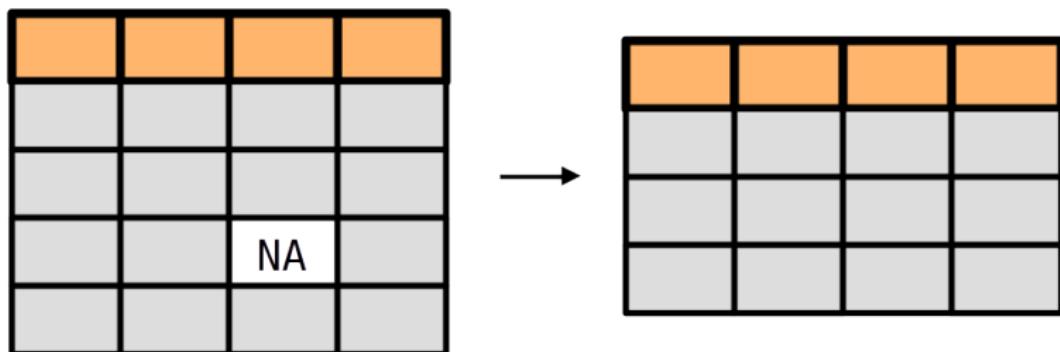
`arrange()` sorts rows.



`filter()` reduces to specific rows.



`drop_na()` removes rows with missing values.



An important skill is to translate ordinary language into dplyr verbs.

Write your own code to respond to five requests.

`summarize()` reduces a column to a single value.

- For example: `sum()`, `mean()`

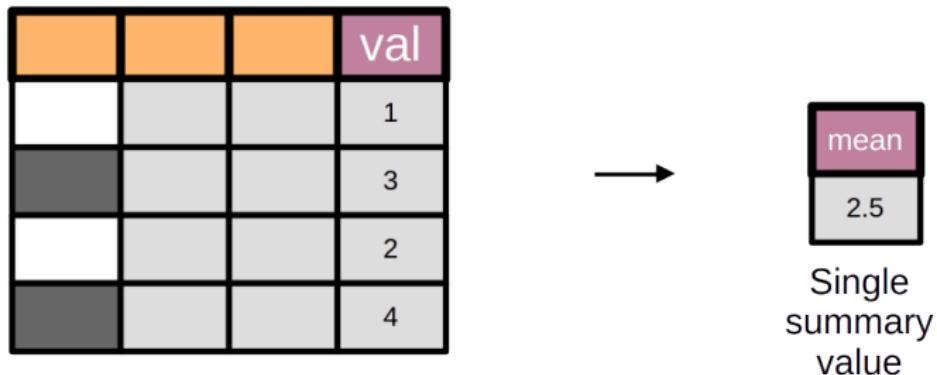
Often combined with `group_by()`.

`group_by()` sets the “grouping” property of the df.

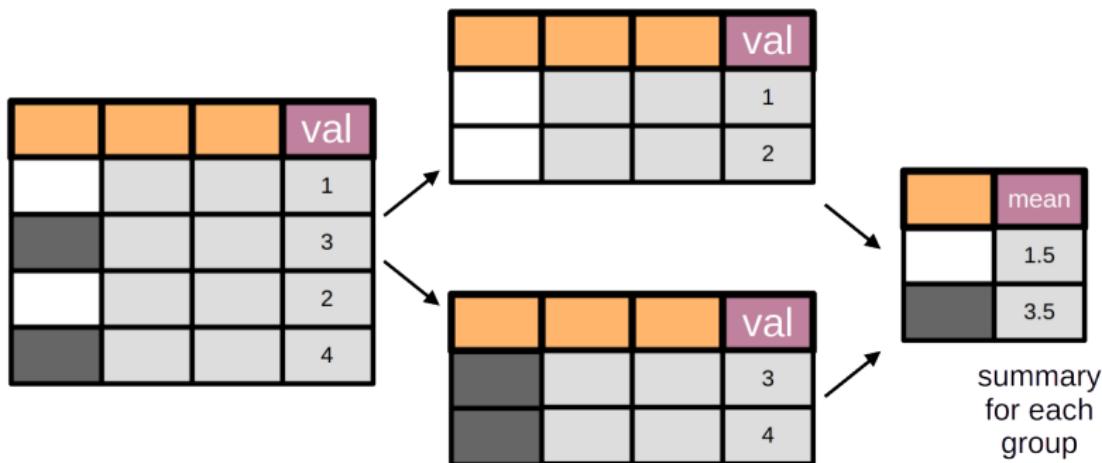
- Gives instructions for future commands
- Changes “level” of calculation

Let's find the mean, min, and max body mass for the three species.

summarize():



group_by() %>% summarize():



Other summary functions:

- `slice_min()` for the lowest values
- `slice_max()` for the highest values
- `count()` as a shortcut for `n()`

`group_by()` has useful applications to `mutate()`.

- Resulting column value depends on group.

How does each penguin's weight compare to the others in its species?

Remove the grouping instruction with ungroup.

We can perform a grouping operation at first, then go back to looking at the entire df at once.

Let's see how it works.

- Run each code chunk as-is
- Predict what happens if you include `group_by()`
- Uncomment the line and see what happens

We can create a new column with `mutate()`, such as the product of two other columns.

We can create a column based on a condition with `case_when()`.